[Your Name]

M.S. in Computer Science, Zhejiang University | incoming Sept 2026

[email] • GitHub: github.com/[handle] • [homepage / Google Scholar]

RESEARCH INTERESTS

Large language models, with a focus on **post-training and reinforcement learning for reasoning** — reward design, credit assignment, training stability (entropy dynamics), and the pre-training → SFT → RLHF pipeline.

EDUCATION

Zhejiang University — College of Computer Science and Technology Sept 2026 – Jun 2029 (expected)

M.S. in Computer Science and Technology • Advisor: [Prof. Name]

[Undergraduate University] [20XX] – 2026

B.S. in [Major] • GPA [X.XX / 4.0], Rank [X / XX]

PUBLICATIONS & PREPRINTS

1. **[EMNLP 2026 paper title].** [Authors, with your position]. *Under review, EMNLP 2026.*
2. **DivGRPO: Divergence-Guided Asymmetric Credit Assignment for Reasoning RL.** [Authors]. *In preparation, targeting AAAI 2027.*

RESEARCH EXPERIENCE

Divergence-Guided RL for LLM Reasoning (DivGRPO) — *lead* 2025 – present

- Identified a core failure of GRPO: a trajectory's final-outcome reward is broadcast to **all** tokens, mis-crediting correct steps in incorrect trajectories and rewarding lucky steps in correct ones — empirically linked to entropy instability (collapse / explosion) on distilled reasoning models.
- Proposed **DivGRPO**: detect token-level divergence points between correct/incorrect trajectories and apply an **asymmetric, multiplicative** reweighting of the GRPO advantage in pure token space (no extra model, hidden states, or logits) — penalty peaks at the divergence point and decays exponentially, while correct paths are reinforced after divergence.
- Implemented on the **verl** framework; ran controlled, multi-seed experiments across **DeepSeek-R1-Distill-Qwen-1.5B / 7B** and **Qwen3-4B** on 8×A800.
- **Result**: DivGRPO achieves stable low-entropy convergence where GRPO is unstable — at 1.5B it preserves AIME24 acc@12 ≈ 0.23 (vs. GRPO ≈ 0.16 under entropy explosion); at 7B it holds macro accuracy ≈ 0.53 throughout training, avoiding the late-stage degradation GRPO exhibits.

[EMNLP 2026 Project --- short title] 2025 – 2026

- [One line: problem]
- [One line: method]
- [One line: quantified result]

SELECTED PROJECTS

LLM from scratch: pre-training → SFT → RLHF [code]

- Implemented Transformer, tokenizer, and the full pipeline from scratch; pre-trained on [XX] tokens, then SFT and PPO/GRPO-based alignment — building end-to-end intuition for reward design, KL control, and entropy monitoring.

Reproducible RL training & evaluation infrastructure (verl)

- Built a training-to-evaluation pipeline (run directories, logging, artifact archival, auto-resume) for multi-GPU GRPO/DivGRPO experiments, plus a unified math-answer verifier (`math_verify` over `\boxed{}` with equivalence checking) for consistent benchmarking.

TECHNICAL SKILLS

- **Languages**: Python (proficient), [C++ / Java (optional)]
- **Machine Learning**: PyTorch, LLM pre-training / SFT / RLHF (PPO, GRPO), verl; distributed training (FSDP), CUDA environment setup
- **Infrastructure**: Linux, tmux, Git; vLLM / SGLang inference (familiar)
- **Spoken languages**: Chinese (native), English (CET-6; academic reading & writing)

HONORS & AWARDS

- [Scholarship / competition / national-level award — year and rank]
- [...]